

Instantly share code, notes, and snippets.

prashants / lwnfs.c

Created 10 years ago

☆ Star

<> Code  Revisions 1 ☆ Stars 8  Forks 22

Updated lwnfs

 lwnfs.c

```
1  /*
2   * Demonstrate a trivial filesystem using libfs.
3   *
4   * Copyright 2002, 2003 Jonathan Corbet <corbet@lwn.net>
5   * This file may be redistributed under the terms of the GNU GPL.
6   *
7   * Chances are that this code will crash your system, delete your
8   * nethack high scores, and set your disk drives on fire.  You have
9   * been warned.
10  */
11  #include <linux/kernel.h>
12  #include <linux/init.h>
13  #include <linux/module.h>
14  #include <linux/pagemap.h>      /* PAGE_CACHE_SIZE */
15  #include <linux/fs.h>           /* This is where libfs stuff is declared */
16  #include <asm/atomic.h>
17  #include <asm/uaccess.h>        /* copy_to_user */
18
19  /*
20   * Boilerplate stuff.
21   */
22  MODULE_LICENSE("GPL");
23  MODULE_AUTHOR("Jonathan Corbet");
24
25  #define LFS_MAGIC 0x19980122
26
27
28  /*
29   * Anytime we make a file or directory in our filesystem we need to
30   * come up with an inode to represent it internally.  This is
31   * the function that does that job.  All that's really interesting
```

```
32  * is the "mode" parameter, which says whether this is a directory
33  * or file, and gives the permissions.
34  */
35  static struct inode *lfs_make_inode(struct super_block *sb, int mode)
36  {
37      struct inode *ret = new_inode(sb);
38
39      if (ret) {
40          ret->i_mode = mode;
41          ret->i_uid = ret->i_gid = 0;
42          ret->i_blocks = 0;
43          ret->i_atime = ret->i_mtime = ret->i_ctime = CURRENT_TIME;
44      }
45      return ret;
46  }
47
48
49  /*
50   * The operations on our "files".
51   */
52
53  /*
54   * Open a file. All we have to do here is to copy over a
55   * copy of the counter pointer so it's easier to get at.
56   */
57  static int lfs_open(struct inode *inode, struct file *filp)
58  {
59      filp->private_data = inode->i_private;
60      return 0;
61  }
62
63  #define TMP_SIZE 20
64  /*
65   * Read a file. Here we increment and read the counter, then pass it
66   * back to the caller. The increment only happens if the read is done
67   * at the beginning of the file (offset = 0); otherwise we end up counting
68   * by twos.
69   */
70  static ssize_t lfs_read_file(struct file *filp, char *buf,
71                              size_t count, loff_t *offset)
72  {
73      atomic_t *counter = (atomic_t *) filp->private_data;
74      int v, len;
75      char tmp[TMP_SIZE];
76
77      /*
78       * Encode the value, and figure out how much of it we can pass back.
79       */
79      v = atomic_read(counter);
80      if (*offset > 0)
```

```
81         v -= 1; /* the value returned when offset was zero */
82     else
83         atomic_inc(counter);
84     len = snprintf(tmp, TMP_SIZE, "%d\n", v);
85     if (*offset > len)
86         return 0;
87     if (count > len - *offset)
88         count = len - *offset;
89 /*
90  * Copy it back, increment the offset, and we're done.
91  */
92     if (copy_to_user(buf, tmp + *offset, count))
93         return -EFAULT;
94     *offset += count;
95     return count;
96 }
97
98 /*
99  * Write a file.
100  */
101 static ssize_t lfs_write_file(struct file *filp, const char *buf,
102                             size_t count, loff_t *offset)
103 {
104     atomic_t *counter = (atomic_t *) filp->private_data;
105     char tmp[TMP_SIZE];
106 /*
107  * Only write from the beginning.
108  */
109     if (*offset != 0)
110         return -EINVAL;
111 /*
112  * Read the value from the user.
113  */
114     if (count >= TMP_SIZE)
115         return -EINVAL;
116     memset(tmp, 0, TMP_SIZE);
117     if (copy_from_user(tmp, buf, count))
118         return -EFAULT;
119 /*
120  * Store it in the counter and we are done.
121  */
122     atomic_set(counter, simple_strtol(tmp, NULL, 10));
123     return count;
124 }
125
126
127 /*
128  * Now we can put together our file operations structure.
129  */
```

```
130 static struct file_operations lfs_file_ops = {
131     .open    = lfs_open,
132     .read    = lfs_read_file,
133     .write   = lfs_write_file,
134 };
135
136
137 /*
138  * Create a file mapping a name to a counter.
139  */
140 static struct dentry *lfs_create_file (struct super_block *sb,
141     struct dentry *dir, const char *name,
142     atomic_t *counter)
143 {
144     struct dentry *dentry;
145     struct inode *inode;
146     struct qstr qname;
147
148     /*
149      * Make a hashed version of the name to go with the dentry.
150      */
151     qname.name = name;
152     qname.len = strlen (name);
153     qname.hash = full_name_hash(name, qname.len);
154
155     /*
156      * Now we can create our dentry and the inode to go with it.
157      */
158     dentry = d_alloc(dir, &qname);
159     if (! dentry)
160         goto out;
161     inode = lfs_make_inode(sb, S_IFREG | 0644);
162     if (! inode)
163         goto out_dput;
164     inode->i_fop = &lfs_file_ops;
165     inode->i_private = counter;
166
167     /*
168      * Put it all into the dentry cache and we're done.
169      */
170     d_add(dentry, inode);
171     return dentry;
172
173     /*
174      * Then again, maybe it didn't work.
175      */
176     out_dput:
177     dput(dentry);
178     out:
179     return 0;
180 }
```

```
179  /*
180  * Create a directory which can be used to hold files. This code is
181  * almost identical to the "create file" logic, except that we create
182  * the inode with a different mode, and use the libfs "simple" operations.
183  */
184  static struct dentry *lfs_create_dir (struct super_block *sb,
185                                       struct dentry *parent, const char *name)
186  {
187      struct dentry *dentry;
188      struct inode *inode;
189      struct qstr qname;
190
191      qname.name = name;
192      qname.len = strlen (name);
193      qname.hash = full_name_hash(name, qname.len);
194      dentry = d_alloc(parent, &qname);
195      if (! dentry)
196          goto out;
197
198      inode = lfs_make_inode(sb, S_IFDIR | 0644);
199      if (! inode)
200          goto out_dput;
201      inode->i_op = &simple_dir_inode_operations;
202      inode->i_fop = &simple_dir_operations;
203
204      d_add(dentry, inode);
205      return dentry;
206
207  out_dput:
208      dput(dentry);
209  out:
210      return 0;
211  }
212
213
214
215  /*
216  * OK, create the files that we export.
217  */
218  static atomic_t counter, subcounter;
219
220  static void lfs_create_files (struct super_block *sb, struct dentry *root)
221  {
222      struct dentry *subdir;
223
224      /*
225       * One counter in the top-level directory.
226       */
227      atomic_set(&counter, 0);
228      lfs_create_file(sb, root, "counter", &counter);
```

```
228  /*
229  * And one in a subdirectory.
230  */
231      atomic_set(&subcounter, 0);
232      subdir = lfs_create_dir(sb, root, "subdir");
233      if (subdir)
234          lfs_create_file(sb, subdir, "subcounter", &subcounter);
235  }
236
237
238
239  /*
240  * Superblock stuff. This is all boilerplate to give the vfs something
241  * that looks like a filesystem to work with.
242  */
243
244  /*
245  * Our superblock operations, both of which are generic kernel ops
246  * that we don't have to write ourselves.
247  */
248  static struct super_operations lfs_s_ops = {
249      .statfs      = simple_statfs,
250      .drop_inode  = generic_delete_inode,
251  };
252
253  /*
254  * "Fill" a superblock with mundane stuff.
255  */
256  static int lfs_fill_super (struct super_block *sb, void *data, int silent)
257  {
258      struct inode *root;
259      struct dentry *root_dentry;
260  /*
261  * Basic parameters.
262  */
263      sb->s_blocksize = PAGE_CACHE_SIZE;
264      sb->s_blocksize_bits = PAGE_CACHE_SHIFT;
265      sb->s_magic = LFS_MAGIC;
266      sb->s_op = &lfs_s_ops;
267  /*
268  * We need to conjure up an inode to represent the root directory
269  * of this filesystem. Its operations all come from libfs, so we
270  * don't have to mess with actually *doing* things inside this
271  * directory.
272  */
273      root = lfs_make_inode (sb, S_IFDIR | 0755);
274      if (! root)
275          goto out;
276      root->i_op = &simple_dir_inode_operations;
```

```
277     root->i_fop = &simple_dir_operations;
278 /*
279  * Get a dentry to represent the directory in core.
280  */
281     root_dentry = d_alloc_root(root);
282     if (! root_dentry)
283         goto out_input;
284     sb->s_root = root_dentry;
285 /*
286  * Make up the files which will be in this filesystem, and we're done.
287  */
288     lfs_create_files (sb, root_dentry);
289     return 0;
290
291 out_input:
292     input(root);
293 out:
294     return -ENOMEM;
295 }
296
297
298 /*
299  * Stuff to pass in when registering the filesystem.
300  */
301 static struct dentry *lfs_get_super(struct file_system_type *fst,
302                                     int flags, const char *devname, void *data)
303 {
304     return mount_bdev(fst, flags, devname, data, lfs_fill_super);
305 }
306
307 static struct file_system_type lfs_type = {
308     .owner          = THIS_MODULE,
309     .name           = "lwnfs",
310     .mount          = lfs_get_super,
311     .kill_sb        = kill_litter_super,
312 };
313
314
315
316
317 /*
318  * Get things set up.
319  */
320 static int __init lfs_init(void)
321 {
322     return register_filesystem(&lfs_type);
323 }
324
325 static void __exit lfs_exit(void)
```

```
326 | {  
327 |     unregister_filesystem(&lfs_type);  
328 | }  
329 |  
330 | module_init(lfs_init);  
331 | module_exit(lfs_exit);
```

mbaynton commented on Jan 28, 2014

Try my fork if this code produces an error when mounting

RadNi commented on Jan 10, 2019

For kernel version up to 4.15 use my fork.